

Referee's report on "Playing one's part" (ROPP58)

This paper is about joint intending and joint acting and especially about the problem whether an intentional joint action requires a joint or shared intention to explain it. The author regards as untenable both the hypothesis that this requires that the agents intend, severally, that they do the thing in question, and also the hypothesis that they intend, jointly, that they do that thing. His view is rather merely that, jointly, the agents do that thing. This view does not require that a state of intending is jointly had by the agents. He explicates his own view by means of the notion of participation (deciding and intending to participate) in doing a thing jointly with others.

The author makes a distinction between the "adverbial" and the "sentential" use of "jointly" in the context of action and intention. A logical explication, based on earlier work by some logicians, is given on pp. 3–4. Basically, it seems right, although it is left open how to interpret the notions involved in it. In the author's Löwenheim–Skolem example one might possibly view Löwenheim's and Skolem's action tokens, say *l* and *s*, either as a conjunctively forming the only element *l*&*s* in the set *A* (p. 3) instead considering them as two elements as done in the paper. Generally speaking, the author's account seems to depend on how actions in *A* are individuated and how they are connected constitutively or causally defined (nothing is said about this issue in the logical treatment).

The author discusses intentional action and gives the following account of intentional joint action:

Restricted intention hypothesis*: For any agents a_1, \dots, a_n and act-type *X*, if *X* is jointly done by a_1, \dots, a_n and jointly done intentionally, rather than unintentionally, in execution of a prior joint decision of theirs that jointly, they *X*, then this is, at least in part, because a_1, \dots, a_n intend that, jointly, they *X*, and their intention(s) that, jointly, they *X* control of their *X*-ing.

This thesis connects intentional joint action with intending jointly to perform the action in question. This seems to face the danger of being circular, because "remaining jointly decided" can be claimed to entail joint intending. This issue is worthy of clarification.

Joint intending in the restricted intention hypothesis can be of three

kinds according to the author:

(S) a_1, \dots, a_n severally intend that, jointly, they X;

(J) a_1, \dots, a_n jointly intend that, jointly, they X;

(V) a_1, \dots, a_n variously intend that, jointly, they X, i.e. they intend this by some of them jointly intending it, and the others also intending it.

The author focuses on criticizing (S) and (J) because, according to him, only they have been defended in the literature. (However, see Tuomela, *The Philosophy of Sociality*, Oxford UP, 2007, Chapter 5, for a version of (V).)

The main argument of the paper is goes in terms of an alleged counterexample, "the case of reluctant Ollie" (p. 13 ff.), against (S) and (J) and a lengthy discussion following it. This argument is deductive one and is meant to refute these theses. Its conclusion is that an agent participating in intentional joint action need not himself intend to participate in it. I will not here reproduce the argument, although I refer to its numbered clauses.

The author's example does not seem to give a consistent story, for the reasons that the author himself nicely outlines. There are, however, other descriptions of closely related cases that are consistent. These are different stories, and, to put it bluntly, thus they do not serve the author's purpose. I cannot here try to summarize the long and complex discussion, which gets quite complicated and is not very easy to follow. To pick out one topic, let me here only comment on the author's use of the term "ought". It seems that "ought not" is used differently in clauses (ii) and (iv) of the argument (p. 14). In (ii) it means rationally impermissible in the sense of "is not ideal, optimal, maximal, or the best choice" (which does not entail akrasia). In contrast, in (iv) it means rationally impermissible in a sense that is pertinent to akrasia. Given the two different kinds of "oughts" in the argument, the deduction does not go through.

There are some rather compact and difficult passages in the discussion of the main argument. Thus, for instance, on pp. 25–26 odd symbolism is used such as "abR". This is supposed be some kind of relation–expressing sentence, but is not a well–formed

formula in any logic that the present reviewer is familiar with. These kinds of passages simply do not satisfy ordinary or even minimal readability requirements.

It is not possible here to discuss the contents of this rich paper in more detail, and I will just take up one central topic discussed in the paper. It relates to the aforementioned theses about joint intending:

The adverbial interpretation of (J) amounts to "a₁,...,an intend, jointly, that they jointly X", X being an action, while the sentential interpretation gives "Jointly, a₁,...,an intend that, jointly, they X". The author rejects the adverbial interpretation of (J) on the basis of alleged consensus in the literature and says this of the sentential interpretation (p. 20): When read in the sentential way, "(J) does not predicate any token state or attitude of intention of a plurality of agents [contrary to its adverbial interpretation]. "It does not require that any state is a state that several minds are in, still less that any state is a state of a mind, had by several agents. It requires only that some intention is some states, borne by some agents." (p. 20) Clearly, at least no modern author accepts the last alternative concerned with group minds of the kind that can non-metaphorically intend. Instead, a modern version of the adverbial (J) seems quite possible. Thus consider you and me having the joint intention to carry a table. We can here represent this linguistically by you and me endorsing the same statement expressing joint intention:

I: We will do X together;

You: We will do X together.

We have here two token mental states of intending with the same content. We can stipulate that a joint token of the joint intention to do X together requires this together with a doxastic condition like mutual knowledge about our being in these mental states and with some further conditions that relate us, e.g. our dispositions to make relevant inferences and our being collectively committed to satisfying the intention content by participating in our joint performance of X. Common sense examples (e.g. our joint intention to lift the table) support the possibility of this kind of adverbial use of "jointly" in some cases at least. By accepting the participation requirement the agents satisfy either an adverbial or a sentential version of something like the author's (P*). Let me note that actually

the participation condition—qua an intention to perform one's part as one's part of X and being collectively committed to performing X—is to be found in earlier works such as Tuomela and Miller (1988) [the author incorrectly writes 1998 and for Tuomela's 1995 book he wrongly gives the year 2005].

This paper has been hard to read and its intellectual and philosophical merits are hard to evaluate. The discussion is in general of high quality, yet there are highly problematic arguments in the paper. This compact paper is very long relative to the amount of material it purports to handle. The paper could be made much shorter. There is material that is already "in the public domain" and need not be repeated or summarized in the paper. This point applies especially to the section "Intention as an explanans" (which actually is a bit of a misnomer as very little is said about explanation) and to the Kutz section. The discussion of the Ollie example should be more streamlined and much clearer. As the paper presently stands it is not too readable, and I am sure the readers of ROPP would find it very hard to digest it.

I would recommend that the paper be conditionally accepted, the condition being that the author takes into account the critical points made above. Yet, my first choice, based on thinking of both the author's and the readers' "good" in terms of readability and communicability, is to require a more wholesale rewriting and resubmission. I leave it to the editors to make the final choice between these two alternative recommendations.

The paper is interesting, subtle and illuminating. The discussion of the background metaphysics is careful and helpful, and the main proposal that is put forward in the paper is original and worthy of serious consideration. There are, however, certain points that could do with further clarification, as they are particularly relevant to the main line of argument that is offered. These concern the notions of decision, joint decision, and intention in play (and the connections between them).

In the section 'Intention as Explanans': The following proposal is considered (Singular Case*): "For any agent a and act-type φ , if φ is done by a , and done intentionally, rather than unintentionally, then this is, at least in part, because a intends that he φ and his intention that he φ controls his φ -ing." Counterexamples to the proposal are offered. In order to assess the putative counterexamples we need to know more about the conception of 'intending to φ ' in play. E.g. does one intend to φ just in case one has an intention of the kind that Bratman outlines in his planning theory of intention? A weaker notion of 'intending to φ ' may allow that one intends to φ if one acts with the aim of φ -ing, or if φ -ing is one's aim/goal, which would allow that one could intend to φ in a case in which one doesn't satisfy all of the conditions that are required for having a Bratman-style intention to φ . And it's not obvious that the weaker notion of 'intending to φ ' is vulnerable to the counterexamples that are offered.

The author's response to potential counterexamples to singular case* is to concede them and offer a modified proposal (Restricted Singular Case*): "For any agent a and act-type φ , if φ is done by a , and done intentionally, rather than unintentionally, in execution of a prior decision of his that he φ , then this is, at least in part, because a intends that he φ and his intention that he φ controls his φ -ing." One worry with this proposal is the following: if the notion of intending to φ in play is the strong Bratman-style notion of having an intention to φ , then we can ask, what is the difference between deciding to φ and acquiring an intention to φ ? If there is none, then given that the notion of "executing a decision" is explicated in terms of the notion of acting "under the control of that decision" the proposal seems to amount to the close to trivial claim that: For any agent a and act-type φ , if φ is done by a , and done intentionally,

rather than unintentionally, under the control of an acquired intention of his that he φ , then this is, at least in part, because a intends that he φ and his intention that he φ controls his φ -ing.

According to the Restricted Intention Hypothesis* "For any agents a_1, \dots, a_n and act-type φ , if φ is jointly done by a_1, \dots, a_n and jointly done intentionally, rather than unintentionally, in execution of a prior joint decision of theirs that jointly they φ , then this is, at least in part, because a_1, \dots, a_n intend that, jointly, they φ , and their intention(s) that, jointly, they φ control their φ -ing." We are not told much about what is involved in making a "joint decision", and this notion appears to be important to the argument that follows. The idea that agreement is involved in making a joint decision is alluded to, but what isn't really discussed is the relation between a joint decision and intention. For example, what should be said about the connection between the content of a joint decision and the content of the intentions of the agents who have made that joint decision? Can an agent genuinely count as having jointly decided to do something if he doesn't intend to do the thing that he has jointly decided (with others) to do? If the content of the joint decision plays a role in governing the further planning and decision making that the individual agents involved go on to perform, this may give us reason to think that the content of the joint decision is among the contents of the intentions of the individual agents.

If it is true that one can only genuinely count as having made the joint decision to φ if one thereby intends to φ , then it could be argued that one's joint decision to φ is rational only if the intention to φ that one thereby forms is rational; and the intention to φ that one forms in reaching the joint decision to φ is rational if the joint decision to φ is rational (assuming that relevant circumstances do not change). On this view, the particular constraints that affect the rationality of joint decision making will affect, and be inherited by, the intentions that the individual agents who are involved in the joint decision making go on to form. If the individual agents involved in the joint decision are not rationally criticisable for reaching that joint decision, then each cannot be rationally criticisable for having an intention with the content of the joint decision (assuming that relevant circumstances do not change). This line of argument may affect the argument that the author offers for thinking that Ollie does not intend that Stan and he carry the piano. Further clarification of the connection between the content of a joint decision and the contents of the intentions formed by the agents

involved in that decision would, I think, help to combat this line of objection.