

Editorial Manager(tm) for The Review of Philosophy and Psychology
Manuscript Draft

Manuscript Number: ROPP59R1

Title: Socially extended intentions-in-action

Article Type: Joint Action: What is Shared?

Keywords: joint intentional activity; intention-in-action; own-action condition; basic action; Deborah Tollefsen; John Searle

Corresponding Author: Olle Blomberg

Corresponding Author's Institution:

First Author: Olle Blomberg

Order of Authors: Olle Blomberg

Abstract: According to a widely accepted constraint on the content of intentions, here called the exclusivity constraint, one cannot intend to perform another agent's action, even if one might be able to intend that she performs it. While one may be able to intend that one's guest leaves before midnight, one cannot intend to perform her act of leaving. However, Deborah Tollefsen's (2005) account of joint activity requires participants to have intentions-in-action (in John Searle's (1983) sense) that violate this constraint. In this paper, I argue that the exclusivity constraint should not be accepted as an unconditional constraint on the contents of intentions-in-action: one may intend to perform a basic action that belongs both to oneself and to another agent. Based on the phenomenology of tool-use, I first argue that intentions-in-action of one's basic actions may be technologically extended, meaning that their contents are not restricted to concern the agent's bodily movements. In analogy with this, I then argue that the phenomenology of some skillful joint activities supports the idea that one's basic intentions-in-action may be socially extended, in violation of the widely accepted exclusivity constraint. Tollefsen's account is specifically constructed in order to account for the joint activities of infants and toddlers who lack the capacity to think of others as planning agents and grasp their plan-like intentions (a capacity required by Michael Bratman's (1999) influential account of joint activity). At the end of the paper, I raise some doubt concerning to what extent infants and toddlers have socially extended intentions-in-action.

Response to Reviewers: Reviewer #1:

1) I have changed "several philosophical accounts" to "Bratman's influential account" and then provided references to other accounts that also rely on group-intentions. I have also clarified in what sense Kutz thinks that group-intentions are characteristic for some (but not all) joint activities.

2) I concede that my case was not very well grounded against the possibility that young children might have socially extended intentions-in-action. I have therefore made these remarks less strong and I also mention the intriguing possibility the reviewer brings up (quoting a footnote in Tollefsen's paper).

Reviewer #2:

In response, I have elaborated the examples in the paper in a few places, but only very little. I have also, I hope, explained some of the terminology a bit better.

2) "The expression "joint activities where there is little room for personal-level deliberation and planning" is unclear without any examples of what such activities could be. Whereas some lack of clarity might be fine in the abstract, it would certainly be helpful if more information was provided on page 3 where this expression is mentioned in the text for the first time."

I have added two examples to clarify this.

3) - 6) Done.

7) I have removed the paragraph, which indeed wasn't relevant to the argument.

8) I have added a few references to empirical papers or papers that review empirical work regarding the different stages mentioned.

9) I have clarified the idea of intentions-in-action whose contents are not accessed by conscious experience with an example: "What Searle has in mind here, I think, is that intentional actions that are part of an agent's repertoire of habits and skills may be performed while the agent is attending to some other task or is daydreaming. For example, while being engaged in writing this paper, I may grab the cup to my right side and take a sip of coffee without having a conscious experience of doing so."

10) The expression "pace Bratman himself" has been removed and the sentence changed.

11) "On page 15, what does "a certain restriction" refer to?" I have clarified this a bit by the following:

"The question of whether or not to accept the own-body constraint is a matter of what content it makes sense for an intention-in-action to have, given the kind of mental state that an intention-in-action is."

12) Done.

13) On page 19, last sentence, what aspect of Searle's theory of individual intention and action is referred to here? Is this about the causal link between intention-in-action and its content?

No, this referred to Searle's identification of the experience of acting with the content of an intention-in-action. This has been clarified.

BIG THANKS TO BOTH REVIEWERS!

Author: Olle Blomberg

Title: Socially extended intentions-in-action

Affiliation: Department of Philosophy, School of Philosophy, Psychology and Language Sciences,
University of Edinburgh

Address:

Olle Blomberg

Philosophy

School of Philosophy, Psychology and Language Sciences

Dugald Stewart Building

3 Charles Street

George Square

Edinburgh EH8 9AD

Email: olle.blomberg@gmail.com

Phone: +44 75 997 04 563

Acknowledgements: I want to thank my audiences at the Collective Intentionality VII conference and European Society of Philosophy of Psychology 2010 conference, especially Stephen Butterfill and Auke Pols. For comments and helpful suggestions on earlier paper drafts and presentations, thanks to Natalie Gold, Andy Clark, Matt Nudds, and Ashley Taylor. Thanks to Georg Theiner for forcing me to rethink how the idea of socially extended intentions-in-action is (un)related to vehicle and content externalism. Finally, thanks to Orestis Palermos, Jonas Christensen, Eric Kerr and Andy McKinlay, and two anonymous reviewers. I am grateful to Microsoft Research for funding.

Title: Socially extended intentions-in-action

Abstract: According to a widely accepted constraint on the content of intentions, here called *the exclusivity constraint*, one cannot intend to perform another agent's action, even if one might be able to intend *that* she performs it. While one may be able to intend *that* one's guest leaves before midnight, one cannot intend *to* perform her act of leaving. However, Deborah Tollefsen's (2005) account of joint activity requires participants to have intentions-in-action (in John Searle's (1983) sense) that violate this constraint. In this paper, I argue that the exclusivity constraint should not be accepted as an unconditional constraint on the contents of intentions-in-action: one may intend *to* perform a basic action that belongs both to oneself and to another agent. Based on the phenomenology of tool-use, I first argue that intentions-in-action of one's basic actions may be *technologically extended*, meaning that their contents are not restricted to concern the agent's bodily movements. In analogy with this, I then argue that the phenomenology of some skillful joint activities supports the idea that one's basic intentions-in-action may be *socially extended*, in violation of the widely accepted exclusivity constraint. Tollefsen's account is specifically constructed in order to account for the joint activities of infants and toddlers who lack the capacity to think of others as planning agents and grasp their plan-like intentions (a capacity required by Michael Bratman's (1999) influential account of joint activity). At the end of the paper, I raise some doubt concerning whether infants and toddlers can have socially extended intentions-in-action.

Keywords: joint intentional activity, intention-in-action, basic action, own-action

condition, Deborah Tollefsen, John Searle

1 Introduction

According to Michael Bratman's influential philosophical account of joint intentional activity, if you and I are engaged in a small-scale joint intentional activity *J*, this activity must be outcome of our joint intention to *J*. For us, in turn, to have a joint intention to *J*, each of us must intend that we (continue to) *J* (Bratman 1999, 2009a, 2009b). For example, if you and I are tangoing together, each of us must (among other things) have had an intention of the form "I intend that we tango." Following Christopher Kutz, I will call intentions of this form *group-intentions* (2000, p. 21). A group-intention is an intention that belongs to an individual participant, but whose content refers to the joint activity of all participants. Sometimes the presence of overlapping group-intentions is claimed to be a necessary condition for an activity to count as a truly joint one (Bratman 1999, essays 5 and 6; Pettit and Schweikard 2006; Alonso 2009). At other times, it is presented as at least present in a core class of cases (Bratman 1999, essay 8, 2009a, 2009b; Kutz 2000).¹ However, appeals to group-intentions are problematic since the content of such intentions – "that we *J*" – appear

1 Bratman suggest in "I intend that we *J*" (1999, essay 8) that joint intention may be a multiply realizable phenomenon, so perhaps not all forms of joint intention have group-intentions among their building blocks. Kutz (2000) writes: "Group-intentions are ordinary, instrumental individual intentions whose subject is the individual agent and whose object is a collective act or outcome: I intend that we will dance the tango. Clearly some paradigmatic forms collective action incorporate our action as the direct aim." (p. 21) Kutz main point, however, is that there is a wide range of activities where not all participants have such intentions but which nevertheless qualify as joint activities. A discussion of this point is beyond the scope of this paper, however.

to violate a widely accepted constraint on what can be intended by an agent.² The constraint, which I call *the exclusivity constraint*, says that the performance of another agent's action cannot be part of the content of one's intention to do something (Searle 1983; Kutz 2000; Bardsley 2007; Bratman 2009a, 2009b). This constraint is typically taken for granted as logically implied by what Bratman calls the *own-action condition* on something being an intention, which says that one can only intend to perform one's own actions (1999, p. 148, 2009b, p. 156). Assuming that an action can only belong to one agent, the exclusivity constraint follows from the own-action condition.

Bratman argues that his planning conception of intention allows him to be more liberal about what it is intelligible for an agent to intend (see also Pettit and Schweikard 2006, p. 21). In Bratman's words, "the planning conception of intention supports the legitimacy of the appeal to my intention that *we J*." (1999, p. 98) On this conception, an intention is a plan-like mental state that commits the agent to act. The intentions of others can function appropriately in our plans since we are able to reliably predict that another agent will form a certain intention in many situations. This allows you, for example, to intend that I tell you the time by means of intending to ask me what time it is (Bratman 1999, p.155). Simple examples like this show, convincingly I think, that the exclusivity constraint is not a constraint on the contents

² There is also a problem concerning *the role* that group-intentions play in accounts of joint intention and action (see Petersson 2007). If a joint action must be preceded by a joint intention, and a joint intention requires participants to have personal intentions "that we *J*", where *J* is a joint activity, then the account appears to lead to an infinite regress. I will not discuss this problem here.

of plan-intentions.³

However, reliance on this planning conception of intention risks making participation in joint activity too conceptually or cognitively demanding (Tollefsen 2005; Pacherie 2007, p. 166; Pacherie and Dokic 2006, p. 110). Participants must be capable of both *having* group-intentions and of *perceiving* or *inferring* the group-intentions of other participants. But infants and toddlers engage in what appear to be joint intentional activities (social pretend play, for example), despite the fact that they are not sophisticated “mindreaders”, that is, they lack an explicit understanding of others as agents whose behaviour is driven by mental states such as beliefs, desires, and plan-like commitments to act (Wellman et al. 2001). Furthermore, some joint activities that adult participants spontaneously engage in unfold under constraints that prevent the participants from going through the deliberation and planning that is arguably required for a joint intention (in Bratman’s sense) to be formed. For example, you and I may be strangers who show up at a milonga where we spontaneously start dancing tango with each other upon merely making eye contact. Or we might be acquaintances who bump into each other on the street and start up an informal conversation (another joint activity) that moves from topic to topic in an improvised manner.⁴ While it is no doubt possible to assimilate such cases into Bratman’s planning framework (see e.g. 1999, p. 139), I suspect that a more plausible account of such cases is possible outside

³ Perhaps because plan-intentions are not really intentions in any ordinary sense.

⁴ Bratman (2009a, p.150; also 1999, p.130) takes “our having a conversation together” as an example of the kind of “small scale shared intentional agency” that he is interested in. Kutz (2000) also mentions “conversing” as an example of a joint activity (p. 2).

this framework. But in this paper, I focus on the challenge that Bratman's account face in light of joint activities with infant and toddler participants, who lack capacities for planning and recognition of plans in other agents. To accommodate such activities, joint action theorists must either provide an account of joint activity that does not appeal to group-intentions at all, or argue that group-intentions and recognition of group-intentions are possible even without reliance on Bratman's planning framework. By relying on John Searle's account of intention and action (1993), I will in this paper show that the latter option is plausible, and argue that we can have what I call socially extended intentions-in-action. In rare circumstances, agents cannot merely intend *that we J*, but they can perform an action which includes the performance of another agent's action in its Intentional content.⁵ That is, one can intend (us) *to J* in such a way that the own-action condition is met even if the content of the intention violates the exclusivity constraint.

In the next section, I outline Deborah Tollefsen's (2005) analysis of joint action, which is tailored to deal specifically with the joint activities of infants and toddlers. Her analysis follows the structure of Bratman's account of joint intentional activity closely, but the intentions of the participants are Searlean intentions-in-action (Searle

⁵ I will follow John Searle (1983) in writing 'Intentional' with a capital 'I' when I refer to Intentional content in the technical philosophical sense of the satisfaction conditions of mental states. In other words, not only intentions or actions have Intentional content, but so do beliefs, desires, hopes, etc. The Intentional content of a belief is the conditions under which the belief comes out true. The Intentional content of an intention is the conditions under which the action counts as having been successfully performed.

1983) rather than Bratmanian plan-intentions.⁶ She thus takes something very controversial for granted, namely that an agent can have an intention-in-action with the content “something like ‘that we *J*’” (p. 93). (What she here takes for granted is thus what I argue is possible in this paper.) In section 3, I present the exclusivity constraint and the own-action condition in more detail. Section 4 and 5 contain the meat of the paper. In section 4 I first argue that we can have basic intentions-in-action that are “technologically extended”, and whose conditions of satisfaction thus range beyond movements of the agent’s body. In section 5, I argue that intentions-in-action in an analogous way can be “socially extended”, and that as a consequence, the exclusivity constraint should not be accepted as an unconditional constraint on our intentions-in-action. While I raise some doubts concerning whether Tollefsen’s account is a good starting point for a *general* account of young children’s participation in joint activities, her account may nevertheless be fruitful for understanding other types of joint activities where there is little room for deliberation and planning. Such activities will be ones where the participants are highly skilled at performing the joint activity and where they can trust not only their own skill but on the expertise of the other participants on which they rely. They will also typically be ones where the actions of the participants are tightly coupled to each other. Examples that come to mind include that of two professional dancers or ice skaters performing a dance together, or that of two partners making love. Finally, in section 6, I consider and respond to various objections to the idea that intentions-in-action may be socially extended.

⁶ Searle (1990) has an account of what he calls “collective intentions-in-action” that is very different from Tollefsen’s account. I will not discuss his account in this paper (but see footnote 17).

2 Tollefsen's account

According to Bratman's influential account, a joint intentional activity is the result of a joint intention, and a joint intention is simply a pattern of "interlocking" plan-intentions of the participants about which they have common knowledge. For you and I to have a joint intention, the following conditions must be fulfilled (Bratman 1999, p. 121):

We intend to *J* if and only if⁷

- (1) (a) I intend that we *J* and (b) you intend that we *J*.
- (2) I intend that we *J* in accordance with and because of 1a, 1b, and meshing subplans of 1a and 1b; you intend that we *J* in accordance with and because of 1a, 1b, and meshing subplans of 1a and 1b.
- (3) (1) and (2) are common knowledge between us.

For example, for me and you to be making a vacation trip together, this activity must be the outcome of our having (had) a joint intention to make a vacation trip together.⁸

⁷ Note that Bratman elsewhere suggests that joint intention may be a multiple realisable phenomenon (see 1999, essay 8). In effect, the fulfilment of some other package of conditions may also result in two or more people having a joint intention. Each of the conditions in Bratman's account is an insufficient but nonredundant part of an unnecessary but sufficient package of conditions that realises a joint intention.

⁸ Note that when Bratman uses locutions such as "We intend *to J*" or "They have a joint intention *to J*", they should not be taken as attributions of ordinary intentions to perform an action to a

For this in turn to have been the case, each of us must have had a group-intention of the form “I intend that we go on vacation together”. Furthermore, these intentions must be interdependent. I must intend that we go on a vacation together on the condition that you intend the same, and your intention is likewise dependent on mine. Each of us must also intend that our subplans for carrying out our group-intentions should be made consistent (mesh). If I intend that we go on a vacation together that is focused on clubbing and partying, but you intend that we go on a vacation together that is instead focused on quiet walks in the countryside, and this clash of subplans is known to us, then we do not yet have a joint intention to make a vacation trip together. Finally, it must be the case that the fact that we have these group-intentions and that each intend that our subplans mesh must be something we are mutually aware of. This is a paradigmatic example of joint activity that may be illuminated by Bratman’s account. Both the initial formation of the joint intention and its execution involves deliberation and planning, as well as interdependent commitments to act and common knowledge among participants. Bratman’s account is meant to be quite general in scope though. It is supposed to account for the jointness of activities such as two people singing a duet, playing basketball together, engaging in conversation, or dancing tango (Bratman 1999).

As Tollefsen (2005) points out, the second condition (2) can only be fulfilled if you and I are proficient mindreaders, since it is not merely necessary that we have meshing intentions and plans, but that each of us recognise that this is the case.

collective (hence, attributing a joint intention “to make a vacation trip together” two individuals does not violate the exclusivity constraint).

However, children do not appear to develop a full-fledged representational “theory of mind” until between 3 to 5 years of age (Wellman et al. 2001).⁹ Given that what 1-to-3-year-olds are doing together is appropriately described as joint intentional activity, then this seems to be of kind of joint activity in which condition (2) is not satisfied. Tollefsen (2005) calls this *the mutual responsiveness problem*.¹⁰

I assume here that the mutual responsiveness problem is a real problem for Bratman’s account, but it should be noted that this is not beyond controversy. The idea that children become full-fledged mindreaders between 3 and 5 years of age is mainly based on their performance on so-called false-belief tasks, but what is crucial here is really their understanding of other’s plan-like commitments to act rather than their understanding of other’s false beliefs. In addition, some findings strongly indicate that children may have some implicit grasp of the false beliefs of others already at 15 months of age (Onishi and Baillargeon 2005). However, in light of Bratman’s planning account of intention, which identifies intentions with mental states that fulfil certain roles in practical reasoning and deliberation, I take it that the understanding

9 Following Tollefsen (2005, p. 81), I assume that a full-fledged theory of mind includes the following: (i) an understanding of other persons in terms of their thoughts, intentions, and beliefs; (ii) an understanding that other persons’ thoughts, beliefs, and intentions may differ from one’s own; and (iii) an understanding that others have thoughts and beliefs that may not match with the current state of affairs (false beliefs).

10 Tollefsen points out that Bratman’s common knowledge condition also gives rise to a problem, since having common knowledge arguably also requires having a theory of mind. However, a discussion of this problem and Tollefsen’s proposed solution is beyond the scope of this paper.

that is required to avoid the mutual responsiveness problem is of an explicit rather than an implicit sort. Furthermore, there is as far as I know no positive evidence for the idea that infants and toddlers have an understanding of others as planning agents guided by plan-like commitments to act.¹¹

At any rate, children engage in what appear to be joint intentional activities with each other or with care takers before the age of 3. Between the age of 9 and 15 months, they start to engage in coordinated joint visual attention with both care takers and peers (Bakeman and Adamson 1984; Butterworth 2001). Around the same time they also start to participate in playful activities together with others, such as building a block tower for example (Hay 1979; Verba 1994). In their second year, they even start to engage in “shared pretence scenarios” (Rakoczy 2008, p. 507; Harris and Kavanaugh 1993). For example, a shared pretence scenario might start with an adult pretending to pour tea into two cups and pretending to spill some, to which the child might respond appropriately by pretending to wipe the tea away from the spot where the adult pretended to spill it. Such social play appears to be a clear case of a joint intentional activity.

¹¹ Malinda Carpenter (2009) states that 12-month-olds “arguably show evidence of understanding something about others’ intentions or plans for action—the means others have chosen to use to achieve their goals.” (p. 382) She interprets this as showing that they understand something about Bratmanian plan-intentions and that, thus, the mutual responsiveness problem does not really arise. But Carpenter seems use the notion of a ‘plan for action’ in a very broad sense that encompasses low-level control of movement sequences during action. Bratman’s notion of planning is much more narrow though (this point is made by Stephen Butterfill in an unpublished paper, “Joint action and development”).

According to Tollefsen, the mutual responsiveness problem can be avoided if conditions (1) and (2) are reinterpreted so that the intentions they refer to are *intentions-in-action* rather than plan-intentions. ‘Intention-in-action’ is a technical concept introduced by Searle (1983). It refers to the Intentional component of an action, which specifies the (typically) detailed and fine-grained conditions that the action’s movement component must meet in order for the action to be successful.¹² One of these ‘conditions of satisfaction’ is always that the intention-in-action itself is the cause of the action’s movement component. In other words, intentions are causally self-referential. Crucially for the argument in this paper, the content that one accesses in one’s experience of acting is the content of the action’s intention-in-action. But according to Searle, one can have an intention-in-action without having a *conscious* experience of acting (1983, pp. 91-92). What Searle has in mind here, I think, is that one may automatically perform an intentional action while one is focused on some other primary task or is daydreaming. For example, while thinking about how to formulate this sentence, I may scratch my head or take a sip of coffee without paying any attention to performing these actions. Typically, intentions-in-action are expressed verbally in the form “I am A-ing” or “I am doing A” (1983, p. 84). Deliberate actions also have what Searle calls a *prior intention*, but for actions that are not premeditated, this is the whole story. Prior intentions are expressed in the form “I intend to A” or “I will A” and have more coarse-grained contents than

¹²Searle writes that “in any real-life situation the intention-in-action will be much more determinate than the prior intention, it will include not only that the arm goes up but that it goes up in a certain way and at a certain speed, etc.” (1983, p. 93, dashes added by me)

intentions-in-action (1983, p. 93).

Tollefsen's (2005) argues that intentions-in-action are perceptually overt by means of various behavioural cues such as "facial expression, extended hands, [or] expressions of cooperativeness." (p. 93) The conditions of satisfaction of an intention-in-action are manifest in such cues, and children can according Tollefsen "literally see these conditions." (p. 91) If this is true, then it is possible that an intention-in-action version of condition (2) can be satisfied, even when the participants are children who lack full mindreading capacities (p. 93). There is at least ample evidence which shows that young children can distinguish animate agency from other movements, as well as understand that an agent's behaviour is directed towards a specific goal (Tomasello et al. 2005; Gergely & Csibra 2003). The evidence does not necessarily imply that they can perceive Searlean intentions-in-action, but the evidence is at least consistent with this hypothesis. Furthermore, Tollefsen is not alone in interpreting the evidence in this way (see e.g. Pacherie 2000, 2007; Tomasello, Call and Hare 2003). I tentatively accept this interpretation of the evidence.

While Tollefsen's modification of (2) is a promising solution to the mutual responsiveness problem, I have pointed out that her modification puts the account in a real (rather than merely apparent) conflict with the exclusivity constraint.¹³ While

¹³Insofar as one closely links Searlean intentions-in-action with mechanisms of "motor imagery", then Elizabeth Pacherie's proposed account of joint action in (2007) seems to be very similar to Tollefsen's. Pacherie (2000) herself argues that empirical work on motor imagery throws light on Searle's notion.

children may *perceive* intentions-in-action with a content “something like ‘that we *J*’” (p. 93), how can they or their care takers *have* intentions-in-action with such contents?

¹⁴ After all, such contents are ruled out by the exclusivity constraint. While Tollefsen loosely refers to the content as “something like ‘*that* we *J*’”, intentions-in-action are intentions *to* perform an action, not intentions *that* some state of affairs comes about.¹⁵ (From now on, I will therefore refer to the content of these intentions-in-action as “to *J*” rather than as “that we *J*”.)

3 The own-action condition and two constraints on the content of intentions

Nicholas Bardsley refers to what I call the exclusivity constraint as “an uncontroversial constraint [...] that an individual’s intentions cannot be said to range over others’ actions”, and he uses it to rule out analyses of joint intention as inadequate (2007, p. 144)¹⁶ This constraint is indeed widely accepted and uncontroversial when it comes to what one can intend *to* perform. While one may be able to intend *that* one’s guest leaves before midnight, one cannot intend to perform her act of leaving (Bratman 1999, pp. 115-116, 2009b, p. 157; Pettit and Schweikard 2006, p. 21). Bratman accepts that “what one attempts are [only] one’s own actions.” (1999, p. 97). The contents of what one can intend *to do* are constrained by the own-

¹⁴There is one brief reference on to the possibility that such contents may be problematic: “Aside from the difficulties with the notion of an individual intending that we *J* [...]” (Tollefsen 2005, p. 93) She does not spell out what these difficulties are though, let alone how they could be overcome.

¹⁵Perhaps this is why Tollefsen qualifies the content specification with “something like”.

¹⁶It is clear from the context that Bardsley is primarily referring to “intention-in-action” here.

action condition on something being an intention, according to which “the *subject* of an intending is always the intended *agent* of the intended activity.” (2009b, p. 156) Group-intentions, then, should be glossed as “I intend that [we bring about some state of affairs]”, rather than as “I intend us to [perform some action]”. Accordingly, I take it that Bratman accepts the exclusivity constraint as an unconditional constraint on what one may intend to do.

Other joint action theorists also seem to accept the exclusivity constraint as a constraint on what one can intend to do, among them Kutz (2000, pp. 21-22) and Searle (1983, p. 110). Kutz calls group-intentions “non-standard intentions” because they are not tightly linked to actions and “cannot be directly transformed from intentions that P to intentions to P.” (p. 21) He even suggests that group-intentions might best be thought of as “either expressing a hope that we will do something, or standing in for an individual intention to promote our doing something.” (p. 22) In effect, Kutz excludes the possibility that intentions to P may have a content similar to group-intentions, presumably because he accepts the exclusivity constraint.

Put in terms of Searle’s notion of intention-in-action, the own-action condition says that the content of one’s intention-in-action can only concern the movement or event that together with the intention-in-action itself constitutes one’s action. This movement or event of one’s own action is typically assumed to be *bodily* movements. It is therefore helpful to introduce another possible constraint on the content of intentions-in-action, which I will call *the own-body constraint*. According to this constraint, the content of an agent’s intentions-in-action only concern the agent’s own

bodily movements. I will argue later that if the own-action condition is taken to imply the own-body constraint, then the condition will often be breached in cases where humans perform actions with tools. Searle at least clearly embraces the exclusivity constraint in some form.¹⁷ He claims that if there is “intervening Intentionality” in the causal chain between the tokening of an intention-in-action and the event that is the intention’s outcome, then that action cannot be successful. Searle illustrates this with the following example:

Thus, suppose that unknown to me my arm is rigged up so that whenever I try to raise it, somebody else causes it to go up, then the action is his not mine, even though I had the intention-in-action of raising my arm and in some sense that intention caused my arm to go up. [...] And that this is the right way to construe intentions-in-action is at least indicated by the fact that, when my intentions-in-action make explicit reference to the intentions of other agents, then in general the actions become the actions of those agents. Thus, suppose I know how my arm is rigged up and I want it to go up. My intention-in-action then is *getting the other agent to raise it*, not *raising it*. My action is getting him to raise it, *his* is raising it. (1983, p. 110)¹⁸

¹⁷John Searle thinks that individuals can have a special kind of mental attitude, which he calls “collective intention”, which has the form “We intend to A” or “We are A-ing”. But he embraces the exclusivity constraint as a constraint on the contents of personal intentions-in-action of the form “I am A-ing”.

¹⁸Searle simply writes “intention in action”, without the dashes. I have added the dashes to the quote – so that it reads “intention-in-action” – to emphasise that this is a technical concept of Searle’s.

The latter idea, then, is the following: In general, if I have an intention-in-action that includes in its conditions of satisfaction that it causes you to have an intention-in-action that in turn causes an action *A*, then *A* is performed by you. This excludes me from also performing *A*. The action *I* perform is instead getting you to perform *A*.

Searle thinks that this line of reasoning remains valid even when I am unaware of the mediating role of your intention-in-action. However, it is clearly not valid in all such cases. We can see this by considering one of Searle's own examples (see 1983, p. 98): Gavrilo Princip's blow against Austria and avenge of Serbia (performed by means of his killing of the Archduke Franz Ferdinand in Sarajevo on June 28, 1914). Insofar as "striking a blow against Austria" is an action at all, then it is Gavrilo's intentional action in spite of the fact that most of the links in the causal chain from intention-in-action to massive societal effects must have consisted of "intervening Intentionality". Perhaps Gavrilo was in some sense aware of this intervening Intentionality, at least upon reflection. In that case, Searle's line of reasoning is not even valid for some cases where an agent is aware of the mediating role of others' intentions in the causal chain between intention-in-action and outcome. Given Searle's choice of example in the quote above (raising one's arm), it is reasonable to take the exclusivity constraint to hold only, or only hold unconditionally at least, in cases where *A* is a *basic action*. A basic action is an action *A* that an agent can intend to perform "without intending to do another action by means of which he intends to do *A*." (Searle 1983, p. 100) In most circumstances, squeezing the trigger of a gun is a basic action, since one does not do this by means of intending to contract one's muscles in the arm and hand in such a way that the trigger gets pulled. This is unlike the case when Gavrilo intends to

strike a blow against Austria. He intends to do this by means of firing a gun against the Archduke. In the case of basic actions, it is furthermore hard not to accept the exclusivity constraint as a fundamental constraint. But I will argue in the next section that the constraint does not always hold, even if its scope is restricted to intentions-in-action of basic actions.

To sum up, Tollefsen's modification of Bratman's account leads to a problem. By relying on the notion of intentions-in-action rather than plan-intentions, Tollefsen saws off the branch that allows the original account to avoid violating the exclusivity constraint. Tollefsen does not explain how individuals can have intentions-in-action of the form "I am intending-in-action to J ". I will now suggest a way in which this can be explained.

4 Technologically extended intentions-in-action

My strategy will be to first argue that the conditions of satisfaction of the intention-in-action of a basic action are not always restricted to only concern an agent's own bodily movements. Human agents can have what I call basic but *technologically extended* intentions-in-action. (An intention-in-action is "basic" if it is the intention-in-action of a basic action.) After establishing this possibility, I argue in the next section that basic intentions-in-action can also, contrary to the exclusivity constraint, be *socially extended*. In effect, I am arguing that what may make up "one's own actions" is a lot less restricted than action theorists tend to assume. The own-action condition implies neither the own-body constraint, nor the exclusivity constraint. The reach of a subject's basic actions is not necessarily limited by the subject's body

surface, nor by other agent's co-ownership of the action's physical movement or event.

As far as I can see, there is nothing in Searle's characterisation of intention-in-action as such that restricts the conditions of satisfaction to range over the agent's bodily movements. I here seem to be in disagreement with Joëlle Proust (2003) who interprets Searle as embracing the own-body constraint. She claims that "an intention-in-action fails to represent any sort of further goal, such as 'switching on the light,' 'breaking a vase,' and so on. What it determines is, rather, a bodily movement." (p. 106) It is true that Searle frequently writes that an intention-in-action causes (and presents itself as being the cause of) "a bodily movement", but I think this is a reflection of the choice of examples that he illustrates his theory with, such as "raising one's arm", rather than the expression of a general constraint that is supposed to be part of his theory. Indeed, he says explicitly that we can make "intentional bodily movements where the conditions of satisfaction of our intentions go beyond the bodily movements." (1983, p. 99) However, Searle calls the intentions of such actions "complex intentions". He characterises the self-reflexive content of the intention-in-action of a man who fires a gun in the following way: "This intention-in-action causes it to be the case that the trigger pulls, which causes it to be the case that the gun fires." (1990, p. 409) While the extended "reach" of the action is to be found in the content of the intention, it is represented *as* extended, as pertaining to events beyond the agent immediate control. But there need be nothing complex (in terms of content) about intentions-in-action that reach beyond movements available to the agent considered as an unaided body.

Consider the well-known example of the blind man and his cane, who explores his surroundings by moving and tapping the white cane in front of him (Merleau-Ponty [1945] 2002, pp. 165-166). What, we can ask, is the blind man's 'experience of acting'? His experience is not that of manipulating his arm, wrist and fingers in order to cause the cane to move in certain ways, as if he was holding a cane for the first time in his life. Rather, the blind man simply and directly taps the ground with the tip of the cane. The stick itself has faded out of his awareness and attention, just like our arms and hands are outside the focus of awareness and attention when we "use them" in action. We do not intend to move our joints, wrist and fingers in order to cause our hand to reach out and grasp an object, instead we just reach out and grasp it. Similar things can be said about the perception involved in the skilled use of tools. The blind man's perceptual experience is that of touching the ground at the tip of the cane: he touches the ground *through* the cane. The experience is not that of touching the cane in his hand in order to *infer* what the ground is like. Such extension of our capacities for perception and action is typical of fluent use of tools, where the tools become "transparent" to the user. Following Andy Clark (2008, p. 31), I think that the best picture of the agency involved in such cases is that "of an extended or enhanced agent confronting the (wider) world" rather than that of a bare biological agent facing a tool. Interestingly, neuroscientific findings suggest a subpersonal basis for the shift from phenomenological opaqueness to transparency. It has been shown that an agent's so-called body schema – a neural representation of the body shape and posture – changes as the agent become fluent in using a tool to perform a task (Maravita & Iriki 2004; Clark 2008, sect. 2.5). After learning to use rakes to reach for and collect food,

bimodal neurons in the body schema of Japanese macaques, which normally only fire when the macaque touches something with its hand or sees something near its hand (hence *bimodal*), start to fire when the macaque touches something with the rake or sees something in the space around the rake. Thanks to the plasticity of the body schema, tools become incorporated into the suit of resources at hand that the cognitive system simply takes for granted. These taken-for-granted resources defines the agent's proximal action space. This neuroscientific finding, and the phenomenon of transparency itself, suggests that the intention-in-action that is being carried out in transparent tool-use is not complex, but rather basic and technologically extended. At least within the framework of Searle's theory of intention and action, the phenomenology of action should be taken seriously. After all, the content of the intention-in-action is the content that is accessed when one has the experience of acting.

Searle makes some observations that fit this take on transparent tool-use. First, he notes that what counts as a basic action is relative to an agent's skills (1983, p. 100). Hence, a skilled tool user can act on the world through the tool without having to do this by means of intending to manipulate the tool. Secondly, Searle does not want to restrict the kinds of causal chains that can be involved in the execution of an intention-in-action to those occurring within the body (see 1983, p. 110). If I have an intention-in-action to raise my arm that causes my arm to raise, then this counts as my intention-in-action being carried out successfully, even if someone, unbeknownst to me, has rewired my brain so that the motor commands from my brain go half way around the world via the internet before reaching the effectors. In effect, there is

nothing in Searle's account that precludes that the causal chain between an intention-in-action and the movement it represents extends beyond the agent's body. Besides Searle's habit of calling the intentions-in-action of tool-using actions "complex", his account does not seem to preclude the possibility of basic but technologically extended intentions-in-action.

I conclude that the contents of an agent's basic intentions-in-action are not constrained by the own-body constraint. I do not claim that the own-body constraint on the content's of basic actions is explicitly accepted by many action theorists, but the idea that basic actions are bodily movements seems to be a tacit unexamined assumption in much philosophy of action. For example, Davidson claims "that all primitive actions are bodily movements", where 'bodily movements' is "openhanded enough to encompass such 'movements' as standing fast, and mental acts like deciding and computing." (1980, p. 49) This assumption may be one source of resistance among action theorists to the possibility of intentions-in-action "to *J*". However, at least within the framework of Searle's theory of intention and action, the assumption should be rejected.¹⁹

Since the phenomenon of 'transparency in use' has been used to argue for the view that the vehicles of the contents of perception and action extend beyond the biological boundary of the human (or macaque) organism (Clark 2008, ch. 2; Rupert 2009, ch.

¹⁹For a much more thorough discussion about basic actions and the use of tools, see chapter 3 of

Auke Pols forthcoming doctoral dissertation *Acting with Artefacts* (Eindhoven University of Technology).

8), it is worth pointing out that widening what count as permissible contents of an agent's basic intentions-in-action does not depend on accepting such "vehicle externalism". The question of whether or not to accept the own-body constraint is a matter of what content it makes sense for an intention-in-action to have, given the kind of mental state that an intention-in-action is. It is not a matter of whether the vehicle of the intention-in-action may extend beyond the body surface of the agent. Neither is it a matter of whether the environment of the agent plays a role in "fixing" the content of its intentions-in-action. In other words, rejecting the own-body constraint does not depend on accepting either vehicle externalism or content externalism. The same can be said with regard to rejecting the exclusivity constraint.

5 Socially extended intentions-in-action

In order to reject the exclusivity constraint as an unconditional constraint on what one can intend to do, I will appeal to considerations similar to those that make rejection of the own-body constraint plausible. In other words, I will argue that intentions-in-action can be socially extended, not only technologically extended.

What motivated my notion of basic but technologically extended intentions-in-action were the phenomenon of transparency in skilled tool use. Do people experience a similar shift from opaqueness to transparency with increased participation in a joint activity? I think this sometimes happens. Consider the example of two skilled ice skaters performing a figure dance. Axel Seemann (2009) characterises, from an adopted first-person perspective, their sense of joint control in the following way:

[I]t isn't that *I* experience myself as being in a position to determine *your* doings by my actions. It is, rather that the experience really is one of *us* controlling our doings. And this experience is an embodied one: in the most obvious kind of case, such as the example of the figure dancers, our bodies really seem to form an experiential unit. The awareness I enjoy of your body in a dance is quite unlike the sensation of your flesh pressing against mine that I might be exposed to in a crowded subway carriage. Your body seems to form part of our joint interface with the world. (2009, p. 508)

To my ears, Seemann's observation here rings true. Like in the case of skilled tool use, agents sometimes experience each other as transparent extensions of their own activity. The coordination of actions recedes to the background of attention and the agents experience themselves acting as one attending to their joint activity. The intentions-in-action here are basic, they are not executed via the execution of other intentions, whether one's own or those of one's co-agent.

My claim here is not that the whole activity – the entire dance on the ice for example – is part of the content of the participants' basic intentions-in-action, but rather a move or turn within that activity. My claim is that something like the following joint move may be part of each intention's content: As one of the dancing ice skaters trace a curve on the ice, he lifts his partner up into a position where she has one leg held above her head and the other leg stretched out parallel to the ice (a so-called full Biellmann position). Here, the intention-in-action of each of the skaters will be dependent on the intention-in-action of the other, the move as whole will be a basic

action for each of them. The bodily movements of their partner will be part of the conditions of satisfaction of the intentions-in-action of each of them. Furthermore, the self-reflexive condition that an intention-in-action causes the movement or event that is specified in its content is satisfied for both of them. The intention-in-action of each skater cause not just their own bodily movements, but also the bodily movements of the other. After all, in the absence of their intention-in-action, the movements of the other would not be performed.

I have argued that the notion of ‘socially extended intentions-in-action’ is intelligible. It is also a useful notion, since it captures the the kind of intentions-in-action that participants engaged in some skillful joint activities have, such as figure skaters or dancers. One mark of such joint intentional activities would then be that the participants have socially extended intentions-in-actions that overlap. This is one way in which we might understand the idea that intentions-in-action could have the content “to *J*” that is part of Tollefsen’s account. It is also an alternative way of elaborating Searle’s analysis of intention-in-action in *Intentionality* (1983) to deal with joint intentional activity, which does not appeal to a special “we”-mode of intention-in-action (as in Searle 1990).

6 Objections and rejoinders

One initial objection to socially extended intentions-in-action might be the following. The sources of knowledge about our own and others’ actions are different in important ways. We have no *proprioceptive* information about the bodily posture or movements of others. To the extent that such information is required for carrying out

intentions-in-action and for monitoring the dynamic unfolding of action, there is reason to doubt that intentions-in-action can be socially extended. But no such proprioceptive information is available in tool use either, so such information does not seem to set a limit to what can part of the content of basic intentions-in-action.²⁰

Some evolutionary considerations suggest, however, that the boundary of the biological organism has an epistemic importance which at least makes it unlikely that extended intentions-in-action are common. Even if vision is important for awareness and knowledge of our own actions, it can be argued that proprioceptive information has a kind of privileged status in our subpersonal cognitive machinery that is not enjoyed by perceptual information about the extra-bodily environment. In arguing against the hypothesis of extended cognition, or “vehicle externalism”, Kim Sterelny (2004) draws attention to the fact that an organism is a conglomeration of parts that have evolved together and thus co-adapted to each other over time. Evolution has ensured that the internal “environment” of the organism is a friendly one: Information exchange is reliable, trustworthy and efficient. “[O]ver evolutionary time”, as Sterelny (2004) puts it, “the internal informational environment of an agent will become more transparent.” In contrast, the external environment is a “shared and sometimes contested space” inhabited by other agents who are potentially out to deceive and manipulate you. Hence, while there is selection pressure on perceptual systems to become more reliable and efficient, it is unlikely that external resources,

²⁰Indeed, in circumstances where proprioceptive and visual information about our own actions are in conflict, our conscious reports about our own bodily movements seem to rely more on vision than proprioception, as for example the ‘rubber hand illusion’ demonstrates (Botvinick & Cohen 1998).

such as tools, public information, and (I would add) other agents become as transparent to us as our internal resources. This is because perceptual systems pick up information from an epistemically hostile environment. But as Clark (2008, p. 103) points out in reply to Sterelny, these evolutionary considerations do not show that technological extension of cognition is impossible, just that we should be sceptical about their prevalence.

Given that Sterelny is correct in characterising internal information flow as basically trustworthy and reliable, and external information flow as always potentially noisy and deceptive, technologically extended intentions-in-action ought to be quite rare, and arguably, socially extended intentions-in-action ought to be even more rare. The possibility of defection and deception is much greater in the case of other agents than in the case of tools. The objection, then, is that a socially extended intention-in-action will depend on the absence of a kind of vigilance and double-checking that we should expect to be the norm in interaction with other agents. However, like in the case of technological extension, these considerations do not show that socially extended intentions-in-action are impossible, just that that we should be sceptical about their prevalence. In addition, the epistemic difference between the internal and the external environment shouldn't be exaggerated.²¹ If, as Sterelny himself believes (see 2003), group selection has played an important role in the evolution of our pervasive

²¹In a recent paper, Sterelny (2010) comments on his earlier argument against the extended cognition hypothesis. He submits that his “initial contrast [between the internal and external environment] was overdrawn” and that his “earlier ideas on the importance of contested space were overinfluenced by Machiavellian models of social interaction.” (p. 474)

tendency to cooperate, then human groups are in a sense made up of co-adapted parts that have evolved together. Less controversially, infants and toddlers are profoundly dependent on their caretakers and the group into which they are born. Even if there are cognitive mechanisms in older children and adults with the function of double-checking and vetting during joint activity, it does not seem plausible that such mechanisms are operating in young children, at least not during interaction with close kin or other familiar individuals. Furthermore, my argument crucially depends on the phenomenology of participation in joint activity. In cases where an agent experience his participation in a joint activity as transparent, then that should be reflected in the content of his intention-in-action. I am assuming here that Searle is right to identify an agent's experience of acting with the content of the agent's intention-in-action: The content of one's experience of acting just *is* the content of one's intention-in-action.

However, perhaps the experience associated with a socially extended intention-in-action is always mistaken or illusory. One might grant that a basic intention-in-action may have conditions of satisfaction that makes it socially extended, but argue that, in fact, the self-reflexive condition that the intention-in-action causes a jointly performed movement or event will never be satisfied. This would account for the phenomenological observations that I have been swayed by, but would arguably make socially extended intentions-in-action less interesting since they would not be appropriately causally connected to the (joint) activity. It is not clear why one should think that the self-reflexive condition cannot sometimes be satisfied though. After all, an intention-in-action is just supposed to be one in a host of necessary factors that bring a movement or event about. The agent's body and the world must cooperate for

the right things to happen even in the case of ordinary solo action.

One could object that socially extended intentions-in-action cannot have a *joint* or *shared* activity as its content, as required by Tollefsen's account. I have argued that the bounds of what counts as one's own basic actions are wider than action theorists have thought, but this does not make those actions joint or shared. Hence, the objector continues, I have not here gone to the heart of what is philosophically puzzling about joint intentional activity, namely the possibility of whether to do this or that being up to us, a matter for us to settle together. As David Velleman (1997) points out, I can only intend to do something which is up to me to settle. But if it is up to me to settle whether or not to take a course of action, then it cannot simultaneously be up to you to settle whether to do this or not. The puzzle is how an agent to exercise control over something while at the same time delegate this control to someone else.

What I am proposing here is not a solution to this puzzle. I am not making a proposal about how agents can share a future-directed intention or decision to engage in a joint intentional activity, I am rather trying to suggest what characterises some joint intentional activities, irrespectively of whether the participants jointly chose to enter the activity. I take it that many joint activities that kids engage in, for example, are characterised by what Bratman calls "mutual responsiveness in action", but not necessarily "mutual responsiveness of intentions" (see 1999, essay 5). What is joint is control and guidance of the unfolding of the joint activity, not a choice about whether or not to initiate it. Furthermore, I am not saying that the presence of overlapping socially extended intentions-in-action is the only condition required for an activity to

be jointly intentional. I have merely focused on Tollefsen's first condition, but for two participants to engage in joint action, conditions (2) and (3) must also be met. The jointness of joint action must be captured by the analysis as a whole, not by the each participant's intention-in-action "to *J*".

That it is possible for participants' intentions-in-action to be socially extended does not, of course, establish that the intentions-in-action of infants and toddlers can be so extended. Indeed, infants and toddlers do not appear to be in a good position to establish the kind of fluent expertise that makes coordination during joint activity transparent. For this reason, I doubt that an account based on socially extended intentions-in-action, which I suggest that Tollefsen's account is, is applicable to children's joint activities *in general*, including such cases such as two toddlers building a block tower together or engaging each other in social pretend play. But perhaps the conditions needed for two children, or a child and a care-taker, to have overlapping socially extended intentions-in-action do obtain in some simple forms of social play that involve repetition and predictable turn-taking.²² This seems especially plausible in asymmetrical joint activities where a child is playing with an adult, since an adult co-participant is typically more disposed to be controlled and "manipulated" by the child than a peer will be (the adult can thus make up for the child's lack of

²²One of the reviewers suggested that breastfeeding might be an example of a joint activity where a child and a care taker have meshing socially extended intentions-in-action. This, I agree, looks like a plausible case. Of course, breastfeeding will often occur before the participants are able to fulfill conditions (2) and (3) in Tollefsen's account, but in many cultures, breastfeeding continues well beyond the child's first year, in which case all three conditions may indeed be met.

fluent expertise). In a footnote, Tollefsen even raise the intriguing possibility that, in light of very young children's profound dependency on their care takers, their intentions-in-action are actually socially extended as a rule rather than as an exception (2005, p. 93 n. 24).²³ Given the role of phenomenology in my argument for the possibility of socially extended intentions-in-action, and the difficulty of accessing what infants and toddlers experience, it is hard to draw any firm conclusion about what the role, if any, socially extended intentions-in-action might play in young children's participation in joint activities.²⁴ The conclusion that such intentions-in-action have a role to play in some joint activities with adult participants, however, is much more secure.

Even in the case of adult participants, it should be noted that the kind of sustained experience of fluency described by Seemann is the exception rather than the rule. As Kutz (2000) points out, "the complexity arising from problems of coordination renders collective activity salient, making it stand out against a background of unreflective, self-regarding activity." (p. 12) In line with my reflections on Sterelny's earlier argument against extended cognition, he also notes that the possibility of free-riding makes joint activity unlikely to become automatised and unreflective (p. 13). However, Kutz grants, in line with what I have been saying, that "well-rehearsed joint action may require no conscious deliberation or reflection." (p. 12) Indeed, in the case of some joint actions, such as when the ice dancers perform their well-rehearsed joint

²³This possibility was also raised by one of the anonymous reviewers.

²⁴Thanks to an anonymous reviewer for pointing this out.

move, conscious deliberation and reflection will not even be possible given the speed and timing requirements that must be met to avoid failure (and most likely, injury).

I have focused on arguing that intentions-in-action for basic actions can be socially extended, and while it does not seem plausible that infants and toddlers have such intentions-in-action, why must they be basic? Why cannot infants and toddlers simply have *complex* intentions-in-actions with the content “to *J*”? The content of each participants intention-in-action would then specify that it should cause certain bodily movements in turn causes – in combination with the causal power of the other participant’s bodily movements – the *J*-ing to come about. Perhaps this is possible, but it looks like it might not then qualify as a core case of joint activity. For a paradigmatic case of joint action, it seems right to say that the participants should intend to perform the joint action, not merely that all the participants should intend to perform their bit of it, leading to the effect that *J* is performed if their intentions happen to be aligned with each other in the right way.

7 Conclusions

I pointed out that couching Bratman’s account of joint intention in terms of participants’ intentions-in-action in order to avoid the mutual responsiveness problem, as Tollefsen (2005) does, is problematic. Such an account requires young children to have intentions-in-action with the content “to *J*”, but unlike plan-like intentions “*that we J*”, such intentions really do violate the exclusivity constraint. According to this constraint, one’s intentions-in-action cannot range over the actions of another agent. In this paper, I have argued that this constraint should not be accepted as an

unconditional constraint on the possible contents of intentions-in-action. I did this by first showing that there is nothing in Searle's analysis of intention-in-action per se that precludes the contents of one's basic intentions-in-action to range beyond one's bodily movements. The phenomenon of transparency in the skillful use of tools indicates that intentions-in-action often actually do have such wide-ranging contents. Thus, the own-action condition does not imply what I call the own-body constraint. By analogy with such technologically extended intentions-in-action, phenomenological considerations also suggest that intentions-in-action can be socially extended. This means that the bounds of what counts as one's own actions are larger than what action theorists typically assume. Neither technologically extended, nor socially extended intentions-in-action fail to fulfill the own-action condition. As a result, one should reject the exclusivity constraint. Two agents can both have socially extended intentions-in-action with overlapping contents, so that they both guide and monitor the unfolding of their joint activity.

While I have expressed doubt regarding Tollefsen's account as a plausible *general* account of young children's joint activities, I believe the account fills a lacuna in the philosophy of joint action. Bratman mentions the importance of mutual responsiveness of actions during joint activity, but he does not have much to say about it. Tollefsen's account, supported by my notion of socially extended intentions-in-action is interesting in itself, because it suggest one way in which joint activities that depend on intensive online social interaction may be characterised.

References

- Bakeman, R. & Adamson, L.B. (1984). Coordinating attention to people and objects in mother-infant and peer-infant interaction. *Child Development*, 55(4), 1278-1289.
- Bardsley, N. (2007). On collective intentions: collective action in economics and philosophy. *Synthese*, 157(2), 141-159.
- Botvinick, M. & Cohen, J. (1998). Rubber hands 'feel' touch that eyes see. *Nature*, 391(6669), 756.
- Bratman, M.E. (2009a). Shared Agency. In C. Mantzavinos, ed. *Philosophy of the Social Sciences: Philosophical Theory and Scientific Practice*. Cambridge: Cambridge Scholars Press.
- Bratman, M.E. (2009b). Modest sociality and the distinctiveness of intention. *Philosophical Studies*, 144(1), 149-165.
- Bratman, M.E. (1999). *Faces of Intention*. Cambridge University Press.
- Bratman, M.E. (1992). Shared Cooperative Activity. *The Philosophical Review*, 101(2), 327-341.
- Bratman, M.E. (1993). Shared Intention. *Ethics*, 104(1), 97-113.

Bratman, M.E. (1984). Two Faces of Intention. *The Philosophical Review*, 93(3), 375-405.

Butterworth, G. (2001). Joint visual attention in infancy. In J.G. Bremner & A. Fogel, eds. *Blackwell Handbook of Infant Development*. Oxford: Blackwell, pp. 213-240.

Carpenter, M. (2009). Just how joint is joint action in infancy? *Topics in Cognitive Science*, 1(2), 380-392.

Clark, A. (2008). *Supersizing The Mind: Embodiment, Action, and Cognitive Extension*, Oxford University Press.

Davidson, D. (1980). *Essays on Actions and Events*. Oxford: Clarendon Press.

Gergely, G. & Csibra, G. (2003). Teleological reasoning in infancy: the naïve theory of rational action. *Trends in Cognitive Sciences*, 7(7), 287-292.

Harris, P.L. & Kavanaugh, R.D. (1993). Young children's understanding of pretense. *Monographs of the Society for Research in Child Development*, 58(1) [Serial No. 231], v-92.

Kutz, C. (2000). Acting Together. *Philosophy and Phenomenological Research*, 61(1), 1-31.

Maravita, A. & Iriki, A. (2004). Tools for the body (schema). *Trends in Cognitive Sciences*, 8(2), 79—86.

Merleau-Ponty, M. (1945). *Phénoménologie de la perception*. Paris: Gallimard.
Translated. by C. Smith. (2002). *Phenomenology of Perception*. London: Routledge.

Onishi, K.H. & Baillargeon, R. (2005). Do 15-month-old infants understand false beliefs? *Science*, 308(5719), 255-258.

Pacherie, E. (2007). Is Collective Intentionality Really Primitive? In M. Beaney, C. Penco, & M. Vignolo, eds. *Mental processes: representing and inferring*. Cambridge: Cambridge Scholars Press, pp. 153-175.

Pacherie, E. & Dokic, J. (2006). From mirror neurons to joint actions. *Cognitive Systems Research*, 7(2-3), 101-112.

Pacherie, E. (2000). The Content of Intentions. *Mind & Language*, 15, 400-432.

Petersson, B. (2007). Collectivity and Circularity. *Journal of Philosophy*, 104(3), 138-156.

Pettit, P. & Schweikard, D. (2006). Joint Actions and Group Agents. *Philosophy of the Social Sciences*, 36(1), 18-39.

Proust, J. (2003). Action. In B. Smith, ed., John Searle. Cambridge University Press, pp. 102-127.

Rakoczy, H. (2008). Pretence as Individual and Collective Intentionality. *Mind and Language*, 23(5), 499-517.

Rupert, R.D. (2009). *Cognitive Systems and The Extended Mind*. Oxford University Press.

Searle, J.R. (1990). Collective Intentions and Actions. In Philip R. Cohen, Jerry Morgan, and Martha E. Pollack, ed. *Intentions in Communication*. pp. 401-415.

Searle, J.R. (1983). *Intentionality: An Essay in the Philosophy of Mind*, Cambridge University Press.

Seemann, A. (2009). Joint Agency: Intersubjectivity, Sense of Control, and the Feeling of Trust. *Inquiry*, 52(5), 500-515.

Sterelny, K. (2004). Externalism, Epistemic Artefacts and the Extended Mind. In R. Schantz, ed. *The Externalist Challenge*. Walter de Gruyter, pp. 239-254.

Sterelny, K. (2003). *Thought in a Hostile World: The Evolution of Human Cognition*, Wiley-Blackwell.

Sterelny, K. (2010). Minds: extended or scaffolded? *Phenomenology and the Cognitive Sciences*, 9(4), 468-481.

Tollefsen, D. (2005). Let's Pretend! Children and Joint Action. *Philosophy of the Social Sciences*, 35(1), 75-97.

Tomasello, M., Call, J. & Hare, B. (2003). Chimpanzees understand psychological states-the question is which ones and to what extent. *Trends in Cognitive Sciences*, 7(4), 153-156.

Tomasello, M., Carpenter, M., Call, J. Behne, T. & Moll, H. (2005). Understanding and sharing intentions: The origins of cultural cognition. *Behavioral and Brain Sciences*, 28(5), 675-691.

Velleman, J. (1997). How To Share An Intention. *Philosophy and Phenomenological Research*, 57(1), 29-50.

Verba, M. (1994). The beginnings of collaboration in peer interaction. *Human Development*, 37(3), 125-139.

Wellman, H.M., Cross, D. & Watson, J. (2001). Meta-analysis of theory-of-mind development: The truth about false belief. *Child Development*, 72(3), 655-684.