Editorial Manager(tm) for The Review of Philosophy and Psychology
Manuscript Draft

Manuscript Number: ROPP44

Title: Collective Intentionality, Documents, and Social Reality

Article Type: Joint Action: What is Shared?

Keywords: collective intentionality, cooperation, social ontology, documents

Corresponding Author: Dr Giuliano Torrengo, PhD

Corresponding Author's Institution: University of Torino

First Author: Giuliano Torrengo, PhD

Order of Authors: Giuliano Torrengo, PhD

Abstract: The social world is populated by a great many entities, such as promises, contracts, presidents, money, debts, and financial crises. Many philosophers regard collective behaviour and attitudes as the ground of social reality. According to this standard view, social ontology is at bottom composed of collective intentions and cooperative behaviours, and that holds both for simple cases concerning small groups and complex institutional structures. In this paper, this view is challenged and an alternative approach proposed in which the role of collective intentions and cooperative behaviour is very different in (a) small group cases, where it can be seen as the ground of social roles and obligations, and (b) cases concerning complex institutions, in which documents and records of social acts in general are the ground that determines rights, duties, and all other complex status functions.

Suggested Reviewers:

Title: Collective Intentionality, Documents, and Social Reality

Author: Giuliano Torrengo

Email: giuliano.torrengo@unito.it

Keywords: collective intentionality, cooperation, social ontology, documents

Collective Intentionality, Documents, and Social Reality

Abstract
The social world is populated by a great many entities, such as promises, contracts, presidents, money, debts, and financial crises. Many philosophers regard collective behaviour and attitudes as the ground of social reality. According to this standard view, social ontology is at bottom composed of collective intentions and cooperative behaviours, and that holds both for simple cases concerning small groups and complex institutional structures. In this paper, this view is challenged and an alternative approach proposed in which the role of collective intentions and cooperative behaviour is very different in (a) small group cases, where it can be seen as the ground of social roles and obligations, and (b) cases concerning complex institutions, in which documents and records of social acts in general are the ground that determines rights, duties, and all other complex status functions.

Keywords: collective intentionality, cooperation, social ontology, documents

I. Does social ontology rest on a mistake?

Collective behaviour is a genuine phenomenon of human societies. It is a commonplace that people collaborate in many ordinary activities, from team games to class actions. That is the datum on the status of which psychologists, philosophers, and sociologists disagree. Intentions that are, in some sense, collective and shared are likely to have a role in collective behaviour. Roughly, collective intentional acts are expressed by such locutions as "We will do...", or "we believe..." and the like. There are many open theoretical issues concerning the nature of collective intentions. For instance, is the plural form ("We...") of collective intentionality an irreducible feature or can it be analysed in terms of intentions of a singular form ("I will do...", "I believe...")? What is the subject of collective intentionality? Is it the individual or some sort of collective subject? How does a collective intention result in individual contributions to shared actions?

Even though there is disagreement on how to analyse collective intentionality (and not only with respect to the details), few would deny the importance of collective behaviour for understanding society. In particular, it seems to be a shared assumption that collective intentionality is the *ground* of social reality. Philosophers have recently become interested in the ontology of the social sphere. The social world is populated by a great many entities that are particular to it, such as promises, contracts, presidents, money, debts, and financial crises. The status of social entities is peculiar: on the one hand, they would not exist if there were not people behaving in certain ways and entertaining certain thoughts; but on the other hand, they do not depend entirely on the subject as individual thoughts and way of behaving do. If I form an intention to go to the cinema and then change my mind, my intention to go will come into existence and then disappear correspondingly. However, if I *promise* you that I will go to the cinema with you,

my obligation to go is not something that I can dismiss on my own on a whim. In order to dismiss it, I have to engage in some further social activity, such as providing you with a good reason as to why I cannot go. There are different stances towards this seemingly "objective reality" of social objects, from the more realist to the most deflationist and reductionist. However, social ontologists of various kinds assume that collective intentionality is the ground of social reality, in that (i) it is involved in the ontological dependence of social objects on people, and (ii) it has an *explanatory role* in the analysis of the ontology of the social sphere.

In this paper, I want to challenge this common view and sketch an alternative approach. In particular, I want to suggest that the mistake upon which contemporary social ontology rests is the assumption that social reality is constituted, at bottom, by collective intentions with a *shared* content. For instance, on a commonly accepted account of money, the shared belief that certain pieces of paper enable their bearer to exchange them in return for commodities up to a certain established value is what *makes* those pieces of paper *currency*. In order to make a piece of paper into a 5 Euro note, people have to assign it the function of having a value of 5 Euro; that is, they have to share the intentional content that the *piece of paper has a value of 5 Euros*. This, it is maintained, is a case of collective behaviour, akin to cooperative action, and the explanation given above can be extended to *all* aspects of society. Therefore, collective intentionality is the ground of the wide variety of social entities and their complex interrelationships.

I think the explanatory model based on collective intentions with a shared content is incorrect. It works (to a certain extent) for simple cases. However, it cannot be applied everywhere in the social realm. In the alternative approach that I propose herein, the content of a social act, that is, what determines the fundamental features of a given social entity, does not necessarily coincide with that of a shared collective intention. Rather, the relevant contents are those of documents and records that are produced according to certain specified procedures, and are considered thereby valid and binding. This allows us to reduce the role of "social intentions" in the construction of social reality, and put them in the right perspective. I suggest that this new approach is interesting, not only for scholars of social ontology, but also as a working hypothesis for psychological research.

II. Why collective intentionality won't give you the stock market

The commonly accepted view of social reality is John Searle's theory of status function. Social entities depend on people in the sense that they have properties that they would not have if there were no persons. What makes it the case that a piece of paper can function as a 5 Euro note? It is not its microphysical structure or some other intrinsic property. It is the fact that there is a community of people that recognise it as having precisely that function (although there may be physical constraints on the social function: a 12-ton coin would be very impractical). Ontologically speaking, then, social objects *depend* on people both for their existence and their features. If there were no people, there would not be any society, and if there were not certain behaviours and thoughts (of people) *that* society would not exist. According to the standard theory, we find this sort of

dependence all the way down and all the way up in social reality:

> "Collective mental phenomena of the sort we get in organized societies are themselves dependent on and derived from the mental phenomena of individuals. This same pattern of dependence continues higher as we see that social institutions such as governments and corporations are dependent on and derived from the mental phenomena and behaviour of individual human beings" (Searle 2010: 4)

In a sense, this continuity of dependence on the individual intentions from simple collective behaviour to complex institutions is obvious. Once we have accepted that there is no objective spirit that is the agent of collective action, we have to accept that social phenomena depend on individuals. With respect to such a dependence on individuals, there can be no difference between the informal collaboration among people in small groups and that to be found in much more complex situations, such as the Stock Market. This is so, not only in the sense that the existence of an institution requires human beings, but also in the sense that for it to be the kind thing that it is requires pattern of behaviour and thought *of a certain kind*. However, this trivial reading of Searle's passage is not what is meant by the idea that collective intentionality is the ground of social reality.

The standard approach requires not only (a) that all intentionality is in individual minds or brains, and thus that also collective intentionality is at bottom a psychological phenomenon that concerns the individual (independently from the issue of whether its form is irreducibly plural or not), but also (b) that collective intentionality is "the fundamental building block of all human social *ontology*" (Searle 2010: 43, italics mine). That is to say, it is in virtue of collective intentions that there is a domain of social objects that possess characteristic features and have, in some sense, objective existence. Thus, the standard approach endorses two foundational elements.
The first is what is usually called methodological individualism (Searle 1990: 406):
(MI) All intentionality is in individual brains
The second is what we may call the assumption of continuity
(AC) The same sort of dependence of social phenomena on individual thoughts and behaviour that we find in small groups is to be found also in complex social situations

Most philosophers agree with Searle on both issues, although they may disagree on other points, for instance the sense in which social objects have objective existence, or whether collective intentions are primitively plural in their form or can be reduced to singular ones. Here, I do not wish to question the methodological individualism of the standard view, which I embrace, but the grounding role of collective intentionality for social reality that is elicited by (AC).

Collective intentionality manifests itself in two ways. The first concerns action and is *cooperative behaviour*; the second concerns belief and is *collective recognition*. Cooperative behaviour does not require that a complete intentional content be shared: the only part of the intentional content that has to be shared in order for cooperation to occur is the goal. If we cooperate to play a duet, I may play the violin while you play the piano: our intentional contents will be different, but we will share the same goal (*viz.* playing a duet). Collective recognition

requires a more complete shared content: an attribution of a particular function to something or someone. We have collective recognition when people in a community recognise that something has a certain function, as, for instance, when everybody agree that pieces of paper produced in certain ways have a given monetary value.

The two forms of collective intentionality can work in tandem. Consider a simple case. Four friends decide to go camping for the weekend. They cooperate to assign responsibilities for setting up the tent (two people), gathering wood for the fire (one person), and cooking, (one person). As a result of this cooperative assignment, there is collective recognition of the role that each person has. According to (AC), much more complex situations than this one work basically in the same way, namely through cooperation to accomplish shared goals, and collective recognition of status.

Although cooperation is the fully fledged form of collective intentionality, collective recognition seems to be somewhat more fundamental to social reality, in that it yields the minimal requirement for there to be a social domain. Indeed, collective recognition is the minimal requirement for there to be institutions, which are the more complex forms of social organisations:

"[...]the existence of an institution does not require cooperation but but simply collective acceptance or recognition" (Searle 2010: 58)

Let us then focus on collective recognition and institutions. Given (AC), what is claimed about institutional reality is just an extension of what happens at less complex levels of social organisation and *vice versa*. We say that there is an institution X in a given society when we find institutional facts, to the effect that certain people (and/or objects) have properties that they would lack if there were no X. The thesis is that institutional facts are grounded in collective beliefs: in *every* case in which we have an institution (or a social phenomenon in general), people and objects possess the associated social properties *in virtue of a collective recognition of such properties*.

I have two main objections to this thesis. Firstly, taken literally, it leads to implausible consequences regarding the shared *content* of such grounding beliefs. It is widely implausible to maintain that most of the people (in our society) hold detailed beliefs about the social role or function of most of the objects and persons surrounding them. Still, if the social function is a consequence of a collective attribution, then *only a feature of the shared content can account for a feature of the social entity*. If we acknowledge that the social entity has a property that does not result from the collective attribution (and it is not just grounded on its physical constitution, if any), then we need to revise the standard view somehow, because its generalisation from simple cases fails

.

One may have the feeling that such a critique is too hasty, though. It is true that we cannot extend the standard view from simple cases to more complex ones without some adjustments. However, it is quite straightforward to make the adjustments. The collective intentional contents on which social reality depends according to the standard view are not necessarily wholly *explicit*. Even though it is often convenient to describe these intentional contents in terms of explicit

propositional contents, as attributions of a certain function to a certain entity, their psychological status is not necessarily that of a propositional belief. It is feasible to view them as a set of dispositions and capacities.

Consider the following simple situation. I promise you that I will take you to the cinema tomorrow. The social properties (rights and obligations) that you and I acquire by this social act exist in virtue of our explicit beliefs about them. We both believe that I have an obligation and you have a right. In contrast, in more complex situations, much of the social reality that emerges collective intentionality is left to an implicit *deference*. I do not need to have detailed knowledge of the functions that I attribute implicitly to the chief executive officer of an enterprise in respect to his status as CEO, in so far as *other people* have it and I have a disposition to defer to them for such knowledge. We can call the people with the relevant knowledge concerning a given status function "the experts".

This revised view will not suffice. It is implausible to maintain that each social status requires both specific experts about it (or people otherwise entitled to be deferred to) and specific deferential beliefs by the non experts. It might be claimed that the deferential beliefs that ground social institutions are very *general*, so it is plausible to maintain that nearly everyone has them. Yet this is not an answer, but merely a way of rephrasing the problem. Statuses in complex social institutions are complex sets of functions and have complex consequences for society. How can all those sophisticated bindings, constraints, rights, and obligations be *real* if their *only* purported ontological ground comprises rough and general beliefs? Elaborating an account of how deference in such case works and what it requires looks like a promising way to meet this predicament. However, as it will soon become clear, by so doing we will discover that the standard view has to be drastically reconsidered.

The second objection to (AC) is that it does not seem to leave room for mistakes about the social status of someone or something. Even mistakes on the part of individuals are problematic for the standard view. Consider again the example such that I promise you that I will take you to the cinema tomorrow. Tomorrow, though, you think that the day before I promised you that I would take you to a high-class French restaurant. How are we to straighten things out? We seem to have the intuition that *I am right*: I did promise to take you to the cinema and I did not promise to take you to the restaurant. You are wrong: you have misremembered the promise or you are trying to tease me. However, the standard view predicts a stalemate here. The status of my obligation is unclear, because there is no shared collective intentional content, so nobody is right.

One could counterargue that in so far as the *existence* of rights and obligation is concerned, a stalemate is what we should expect. If you and I are the only ones who have partaken in the act of promising, then as soon as you (or I) forget it or fail to recognise it, the promise is no longer binding and so no longer exists. The situation would be different if many people had heard me promising, and even more different if many people were ready to see to it that we kept the promise. In that case, you would be mistaken because the majority think that you are mistaken; exactly as, in general, you have a certain social status only because the majority thinks that you have it. Therefore, mistakes made by individuals are

accounted for by the standard view.

It is not entirely clear to me how the standard view could provide a principled account of how to distinguish the content of the promise from what is believed about it. Besides, it is not clear how to extend the account of mistakes made by individuals in simple cases to account for such mistakes in more complex situations. However, let us grant here that individual mistakes are possible within the standard framework. What about collective mistakes? An entire community may be mistaken about the social status of a person or an object. Consider a football match between team A and team B. Team A wins the match 3-1, but for some strange reason everybody comes to believe that team B won (perhaps everyone in the stadium suffered a drug-induced flashback to the previous year's game, when team B won 3-1, and then the incorrect result was broadcast nationwide over the television and radio). Would not be everybody be *mistaken*? Yet how can this be the case, if the social status of the match depends only on collective attribution and *everybody* believes that B won? On the standard view, they are *not* mistaken; team B has won. Yet this is wrong. If everybody came to know that A scored three goals while B scored only one, but they believed the reverse happened because their beliefs were drug-induced, everybody would realise that they were *wrong*, and not that the result had changed!

One might defend the standard view by claiming that in the case of a football match, the status of the final result is not entirely a matter of what people believe. It also depends on how many goals have been made by each team (even though, of course, what counts as a goal is, in turn, a matter of what people believe). Thus, people may have false beliefs concerning the historical facts and thereby have collective false beliefs about who won. It is quite plausible to think that there is implicit deference to the relevant historical facts here. However, collective mistakes may involve basically every kind of social status, at least in so far as the content of a social act that brings into existence a social entity is *written* in an external accessible support document. For instance, by signing a contract I (usually) accept a lot of duties and assume a lot of rights regarding which I, just like everybody else, may entertain false beliefs. Yet what establishes what are my duties and rights are not my beliefs, nor the beliefs of the collective I am part of, but the content of the written text of the contract. I defer my knowledge to it, because it contains all the information that determines the status function of the persons and objects to which it pertains. In general, it should always be possible for there to be a discrepancy between what someone (or *everyone*) believes about the social status of an entity *x*, and its actual social status. Yet the standard view is unable to allow for such a distinction, unless it is augmented with deference to things such as historical memories, written documents, and the like.

This latter consideration takes us beyond the standard view to an alternative proposal, which will now be presented.

III. Social Intentions and Documents

According to the standard model of social ontology, the collective recognitions on which social reality is grounded consist of cooperation among individuals. Collective recognitions require a shared intentional content: the assignment of a

status function to something or someone. The complexity of many status functions in our modern societies and the possibility of collective mistakes about the social status of something casts some doubts on the tenability of the standard model. For simple situations of collaborative behaviour among members of small groups the model works, but to account for more complex situations it needs to be revised. I think that the problem with the standard model is deep, and that the revisions that need to be made go far beyond small adjustments. In particular, the requirement that the social status of an entity be determined by the content of a shared collective belief about it almost forces us to see recognition as a form of cooperation. But is it?

Now, leaving aside the fact that the similarities between cooperation and collective recognition are rather shallow, the intentions that are required for there to be complex social entities such as institutions do *not* have the form of the collective recognitions that it is plausible to suppose in simpler cases. Consider the following example of a simple case: I promise you that I will pay you 5 Euros by tomorrow. By this act, which requires a community of at least two persons, a social entity is brought into existence: a promise. There are different philosophical theories about the social existence of the promise, but more or less everybody agree that you and I are affected by the promise in certain ways: I have an obligation to pay you, and you have the right to request 5 Euros from me. In other words, the promise has a binding power.

Yet what makes the promise binding? That the parties jointly agree to keep the promise is a precondition for the existence of the social constraints that arise from the act of promising. However, facts that concern *who* is bound, and *how* and *with respect to what* they are bound, depend on the *content* of the social act that the two parties have performed. The content tells us what has been promised to whom, and thus determines the nature of the bond between the parties. In simple situations the standard view works, because the relevant content usually coincides with that of the intentions of the parties involved in the promise, and those intentions have the form of a recognition of a status function.

However, in more complex situations, the content that contains all the social constraints and socially relevant facts, i.e. the content of the social act, is *not that of the actual intentions of the participants of the social act that gives rise to the social object*. As we have seen, people often do not have, even implicitly, detailed beliefs about many status functions, and they can be mistaken about them. More complex social situations require only that the content of the social act be *accessible* to the parties, and not that it be part of the actual intentions of the parties. That is why the content can be *recorded on an external accessible support document.* In modern societies, the content is often recorded on paper or digital media. If we write down that I promised you that I will pay you 5 Euros by tomorrow and *we agree that the content of this inscription is binding*, even if we both forget what has been promised, there will be a way, *in principle*, to settle the issue. Therefore, it is the content produced in the social act and then recorded (even just as a memory trace in people's heads that can be identified by subsequent neurophysiological examination) that warrants the persistence of the social bond and that establishes the nature of the social constraints at issue. I will defend the thesis that such documents are the ontological ground of social entities
.

The difference between the standard view and the one advanced here is deep. On the standard view, what determines the nature of social entities is a shared intentional content. The standard view grounds the elaborated social statuses that we find in complex social structures in rough and general beliefs about such social statuses. However, in general, people have only rough and general beliefs about complex status functions. In contrast, on the proposed view, the intention that is required for the existence of a status function F (the only "social" intention) is *not* necessarily a collective recognition of the status function F. Rather, the shared intentional content consists in recognizing the recorded content of a social act (a document) as *valid*. Thus, social intentions of complex structures *have to be general*, because we need a *nonspecific* intention to regard a certain document as valid in each case in which the feature of the status function is determined by the *specific* content of a document (and not by the content of a shared intention, as may happen in simpler cases).

On the proposed view, the collective recognition that grounds institutional reality is a form of cooperation, but in a much weaker sense than on the standard view. On the standard view, the social world, composed of social entities, exists only if there are collective intentions with a shared *substantive* content, such as *that this piece of paper is a 5 Euro note*. Even if, strictly speaking, a common goal is not required for collective recognition (and even if collective recognition does not need to take an irreducible plural form), the fact that the shared content has to do with the specific features of the social entity makes collective recognition and cooperation two sides of the same coin. In a sense, a shared assignment of status function is a limiting case of a common goal (Searle 2010: 65-8). On the proposed view, the existence of institutions requires only that the individuals in a collectivity share the general belief that the content of certain documents determines social constraints, rights, duties, and the like. The shared part of the intentional content is general, not only in the sense that it does not contain any specific attribution of social functions, but also in that it usually does not concern any specific document. It is required only that we believe that documents *of a certain kind* are considered to be valid in order for there to be institutions. How the content of such intention may be described in detail is probably at least partly an empirical issue that is probably influenced by the cultural background. I propose the following working hypothesis: the content of the shared belief that grounds institutional and complex social reality is that *records of the content of social acts that are produced according to established procedures constitute the court of appeal in cases of dispute*. Which procedures validate a document is a contextual matter: in certain cases it is enough to appeal to the memory of a third-party witness to validate a content, whereas in other cases it is necessary to appeal to a contract on paper or in electronic format, or the law. The (very rarified) form of cooperation that is required in complex social settings is thus very different from the ordinary cooperation that we find in shared actions. It does not concern common goals, not even in the limiting sense of shared recognition of social status, but only the fulfilment of formal requirements.

IV. Collaboration, Negotiation, and Individual Contributions to Social Reality

Collaborative behaviour of a more substantive kind than collective recognition of the validity of certain procedures to produce documents is a widespread phenomenon in society, and it is very important and valuable in many social contexts. The standard view takes cooperation (especially in its degenerative form of the collective recognition of status function) as the basic building block of social reality. However, as I have argued, while this view may work for simple cases, it cannot be generalised to complex cases.

One advantage of the proposed view is that it seems to be better suited than the standard view to explain noncollaborative social behaviour, such as competitive and individualistic behaviour. An objection to the standard view is that collective recognition of status functions seems to require *approval* of the assignment of status functions. However, of course, we do not always approve of the social status that persons and objects around us have. In defence of the standard view, it may be held that the fact that the construction of the social world is supposed to be carried over by the substantive content of the collective intention does suggest that there is an association between recognition and acceptance. However, I take this defence to misfire. At least in so far as there is some *higher-order goal* that is shared, I do not see any problem in accepting the idea that recognition may come with apathy, distress, or even hatred towards the attributed status. If I am playing chess with you and you are winning, I will recognise that you are winning even if you are humiliating me, in so far as I am willing to play with you. Further, if I am willing to live in a democratic society, I will recognise the social status and power of the elected president even if I strongly disagree with his politics.

Many simple situations and certain complex situations, such as the example just given concerning the results of democratic elections, seem to fare well within the standard view. However, is it really plausible to suppose that a specific higher-order collaborative intention is *always* at work in every case in which people behave noncollaboratively? Take a *prima facie* simple case, such as getting fined for parking. Suppose that I think that the officer has made a mistake in fining me. I am not happy at all to recognise the validity of the fine, but this is what I am doing if I want to protest it (if I thought that it is not valid there would be no point in protesting it). The higher-order collaborative intention here would be probably something like acknowledging the power of the institution of the Police. More generally, in so far as we all have the higher-order *general* intention to recognise the legitimacy of the institutions of the society in which we are living, it is always possible to have recognition without approval and for behaviour to be noncollaborative. Again, we seem to have a tension here between the requirement of having intentions with specific content to ground social reality and the *de facto* generic content of actual higher-order collaborative intentions.

The proposed view is more straightforward than the standard approach on this point. The general intention on whose ground noncollaborative social behaviour can occur does not need to be a higher-order one. Generally speaking, complex situations require only the general first-order intention to consider valid the documents that have been produced according to certain procedures. Consider again the case of the fine. If I think that the officer has made a mistake, I will try my best to have the fine dismissed; that is, I protest the validity of a certain document, the official piece of paper that the officer has issued to me upon which is recorded the fine and the circumstances that warranted it. Such a protest can

only be upheld by consulting other, higher-order, valid documents. Further, this is so only because I (as do others) take it for granted that documents produced according to certain procedures are valid. This is what I am implicitly accepting by living in a complex society, and it is very far from any substantive higher-order collaborative intention.

I am not denying that collaborative intention, first-order or higher-order ones, are sometimes the psychological underpinning of social situations. What I am denying is their fundamental role in social ontology. Not every social entity is grounded in an intention to cooperate to achieve a common goal. In modern complex societies, the content of the collectively accepted documents establishes the features of the social world in which we act, sometimes cooperating with each other, other times acting on our own.

Yet, one may have the feeling that the proposed view is too minimalist. People do not always view institutions and society with such apathy. Even granting that the content of collective recognition is often too simplistic and general to ground sophisticated institutional functions, and that we need documents to flesh out the details of status functions, the content of the intentions *is* relevant for the existence of the corresponding social object or function. If *nobody* cared about a certain social function or *everybody* disagreed with its existence, it is difficult to see how such a function could affect our social life, even when there is a document that describes it in full detail.

This objection barks at the wrong tree. The fact that institutional functions require only acceptance of the validity of the relevant documents is compatible with the fact that we *do care* about the content of those documents. The reason we are interested in the content of valid documents is that they can have, and often do have, important consequences for us. Further, we do not just care about documents upon which we conferred validity by collectively *recognising* them as valid, but also about documents that we can contribute to producing. Indeed, our social life is constituted, not only by passive acceptance of the validity of certain documents, but also by active participation in producing documents according to valid procedures.

Now, one may think that at least in those crucial moments when documents are produced, fully fledged cooperation and collective intentions are required. At least in this sense, then, collective intention and shared substantive contents are fundamental in social ontology. Take the creation of laws, which is a paradigmatic case of the production of documents with wide social relevance. Laws usually are fundamental (parts of) documents that are discussed till a collective *agreement* is reached. Without supposing a collective goal, which in many cases is presumably a higher-order one (at least when we are considering disputing parties), it is difficult to see how collective agreement can be achieved. This seems to be true not only of law making, but also of much more ordinary processes of decision making.

I do not deny that this is what happens in elementary cases: we need to agree that the four of us will spend a weekend camping in order to discuss who will be doing what at the camp. However, the will to find an agreement does not require *substantive* collaboration with respect to a common goal in every case. Once we

agreed on what procedure confers validity on documents, it makes perfect sense to say that an agreement is reached on the ground of a *negotiation* regarding individual (or only partially shared) goals. Whether the content of the validated documents is the outcome of an harmonic cooperation between all parties, or rather the outcome of a negotiation regarding conflicting interests, cannot be settled in advance. Modern societies allows for both cases and the minimal behaviour that is required to acknowledge the complex institutions that constitute modern society is collaborative in only a very degenerative form. Rather than requiring complex forms of substantive agreement, complex societies seem to be based on more formal forms of collective attitude, such as agreement on the procedure of validation of documents.

Sometimes there is substantive collaboration towards achieving a common goal, such as when we want to play a duet and we decide that I will be playing the violin while you will be playing the piano. However, at other times, there is agreement only on the formal methods that it is correct to use to produce valid documents, which will be followed by discussion and negotiation to produce content that will satisfy everyone. Collaboration can be a valuable thing in certain circumstances precisely because it is *not* the basic form of social behaviour to be found in every circumstance. A duet is a wonderful thing, just as society in many cases is not.

References

Bratman M.
— 1992, "Shared Cooperative Activity", *Philosophical Review* 101, 2: 327-41

Ferraris M.
— 2009a, *Documentalità. Perché è necessario lasciar tracce*, Roma-Bari, Laterza
— 2009b, "Social ontology and Documentality", *Archiv für Begriffsgeschichte*, Sonderheft 7, pp. 133-147
— 2009c, "Documentality, or Europe", *The Monist*, 92, n. 2, pp. 286-314;

Gilbert M.
— 1989, *On Social Facts*, London, Routledge

Hudin J.
— 2010, "Can Status Functions be discovered?", *forthcoming* in *The Journal for the Theory of Social Behaviour*.

Koepsell, D. and Laurence S. Moss
— 2003, eds., *John Searle's ideas about Social Reality: Extensions, Criticisms, and Reconstructions*, Malden, Mass., Blackwell

Putnam H.
— 1975, "The meaning of "meaning"", *Mind, Language and Reality*, Cambridge, Cambridge University Press: 215-271

Searle J.
— 1990, "Collective Intentions and Actions" in *Intentions in Communication*, P. Cohen, J. Morgan, and M.E. Pollack, eds. Cambridge, Mass. Bradford Books, MIT Press, 1990.
— 1995, *The Construction of Social Reality*, New York, Free Press.
— 2003, "Reply to Barry Smith" in B. Smith & J. Searle, 2003.
— 2010, Making the Social World: The Structure of Human Civilization, Oxford, OUP

Smith B.
— 2003a, "John Searle: From speech acts to social reality" in B. Smith (ed.) *John Searle*, Cambridge, Cambridge University Press: 1-33.
— 2003b, "The Ontology of Social Reality", in B. Smith & J. Searle, 2003.
— 2003c, (ed.) *John Searle*, Cambridge, Cambridge University Press
— 2006, *ms.*, *Document Acts* (http://ontology.buffalo.edu/document_ontology/)

Smith B. & Searle J.
— 2003, "The Construction of Social Reality – An Exchange", *American Journal of Economics and Sociology*, Vol. 62, No. 1: 283 – 309.

Thomasson A.
— 2002, "Foundations for a Social Ontology", *Protosociology: An International Journal of Interdisciplinare Research*, 18-19.

Toumela R.
— 1995, *The Importance of Us: A Philosophical Study of Basic Social Notions*, Stanford Series in Philosophy, Stanford University Press

Tuomela R. and Miller K.
— 1988, "We-intentions", *Philosophical Studies* 53: 367-89

Varzi A.C.
— 2007, "Il denaro è un'opera d'arte (o quasi)", *Quaderni dell'Associazione per lo Sviluppo degli Studi di Banca e Borsa*, 24: 17-39.